

Time-Series Foundation Models (TSFMs) in Finance

Do TSFMs outperform traditional Time-Series prediction models in the field of Finance?

by

Filip Topic

September 2024

10

Dissertation Supervisor: Ramin Okhrati

Dissertation submitted in part fulfilment of the
Degree of Master of Science in Banking and Digital Finance

Institute of Finance and Technology

University College London

15 **DECLARATION**

I, Filip Topic, hereby declare that the work presented in this dissertation is my own original work. Where information has been derived from other sources, I confirm that this has been clearly and fully identified and acknowledged. No part of this dissertation contains material previously submitted to the examiners of this or any other university,
20 or any material previously submitted for any other assessment.

Word Count: 10,000

Name: Filip Topic

Date: 10 September 2024

CONTENTS

| | | |
|----|---|-----------|
| 25 | List of Figures | 4 |
| | List of Tables | 5 |
| | Acknowledgement | 6 |
| | Abstract | 7 |
| | 1 Introduction | 8 |
| 30 | 1.1 Time-Series prediction | 8 |
| | 1.2 Transformer-based Models | 8 |
| | 1.2.1 Time Series Foundation Models (TSFM) | 8 |
| | 2 Background | 9 |
| | 2.1 The Transformer | 9 |
| 35 | 2.1.1 Encoder | 10 |
| | 2.1.2 Decoder | 12 |
| | 2.1.3 Summary | 13 |
| | 2.2 Transformer-based models | 14 |
| | 2.2.1 NLP Transformer-based models | 14 |
| 40 | 2.2.2 CV (computer vision) Transformer-based models | 14 |
| | 2.2.3 Speech and Audio processing | 14 |
| | 3 Related work / Literature review | 15 |
| | 3.1 Section | 15 |
| | 3.1.1 Sub-Section | 15 |
| 45 | 4 Methodology | 16 |
| | 4.1 Time-series prediction evaluation | 16 |
| | 4.2 Data | 17 |
| | 4.2.1 Data Pre-processing | 18 |
| | 4.2.2 Time-series data characteristics | 18 |
| 50 | 4.3 Experiment Design | 19 |
| | 4.3.1 Benchmark | 19 |

| | | | |
|----|----------|---|-----------|
| | 4.3.2 | Time Series Cross Validation | 19 |
| | 5 | Results | 21 |
| | 5.1 | Section | 21 |
| 55 | 5.1.1 | Sub-Section | 21 |
| | 6 | Discussion | 22 |
| | 6.1 | Section | 22 |
| | 6.1.1 | Sub-Section | 22 |
| | 7 | Conclusion | 23 |
| 60 | 7.1 | Section | 23 |
| | 7.1.1 | Sub-Section | 23 |
| | 8 | Appendix | 24 |
| | 8.1 | Brief History of time-series prediction in Finance..... | 24 |
| | 8.2 | Brief History of modern time-series prediction methods..... | 25 |
| 65 | 8.2.1 | Statistical models | 25 |
| | 8.2.2 | Machine Learning (ML) models..... | 26 |
| | 8.2.3 | Deep Learning models..... | 27 |

LIST OF FIGURES

| | | |
|----|--|----|
| | 2.1 Transformer architecture schema [28] | 10 |
| 70 | 2.2 Scaled Dot product attention [28] | 12 |

LIST OF TABLES

ACKNOWLEDGEMENTS

ABSTRACT

Keywords:

1. INTRODUCTION

1.1 Time-Series prediction

Since the beginning of financial markets, their participants have had the desire to predict the future values of instruments being traded there. And for this purposes, they have used many different techniques - majority of which have fallen short of randomly guessing the direction of the price movement. In the past few decades, many models have emerged which claim excellent capability of time-series prediction. First were the statistical models (such as ARIMA), then the Machine Learning models (such as XGBoost), then Deep Learning models (such as DeepAR).

1.2 Transformer-based Models

The revolution in the field of Machine Learning came in 2017 with the invention of the Transformer architecture (more specifically: the attention mechanism). Transformer based models have taken the whole field of ML by storm, however, their application has mostly been in the sub-field of Natural Language Processing. Recently they started seeing use in the field of time-series prediction (eg. Informer).

Time Series Foundation Models (TSFM)

Time-series foundation models are large transformer-based models which are designed for the purpose of time-series prediction and which have been pre-trained on a large amount of time-series data. This research will explore their use on financial time-series data

2. BACKGROUND

2.1 The Transformer

In 2017, Vaswani et al. [28] introduced a ground-breaking model called the "Transformer". Transformer is a model designed to solve the sequence-to-sequence mapping problem. In a sequence-to-sequence problem we have a sequence which consists of tokens¹ coming from a finite vocabulary² which are in a specific order, and we have an output sequence which is again a sequence of tokens in a specific order. The task of a sequence-to-sequence model is to learn the correct relationship between the input and output sequences so that when the model is presented with an unseen input sequence, it can predict what the correct output sequence is. In essence, sequence-to-sequence model's goal is to learn the "meaning" of the relationship between the input and output sequence. An example of sequence-to-sequence tasks are:

1. Language translation - where an input sequence could be a sentence in our native language and the output sequence would be the correct translation of that sentence in a foreign language.
2. Chatbots - where an input sequence could be a question and the output sequence the correct answer to that question.
3. Time-series forecasting - where an input sequence is a certain time-series and the output sequence is the future values of that time-series.

Before the Transformer, there have been many ML models which attempted to solve these tasks. Sutskever et al. (2014) [25] used multilayer LSTM³ (type of RNN⁴ model) for this task. LSTM [15] and GRU⁵ [8] used to be state-of-the art sequence-to-sequence

¹A token is a multi-dimensional numeric vector representation of an element in the sequence. Reason why sequence of tokens is used as input to a sequence-to-sequence model is because ML models can only "understand" numbers - therefore they can't work with some type of sequences (eg. human language sentences). Reason why tokens are vectors rather than simple numbers is because they are designed to represent the semantic meaning of a real-world element they represent and the semantic meaning of a real-world element could change depending on the context. Therefore different dimensions of a token account for different meaning of that element in different contexts.

²The vocabulary doesn't always need to be finite, as we will see later on.

³LSTM stands for Long Short Term Memory cell

⁴RNN stands for Recurrent Neural Network. It is a type of Artificial Neural Network architecture.

⁵GRU stands for Gated Recurrent Unit. GRU is also a type of RNN.

models, however they have some key issues. They require large memory⁶, their nature is sequential⁷, and they suffer from an information bottleneck due to the fixed size of the hidden state⁸. These problems were addressed by the Transformer.

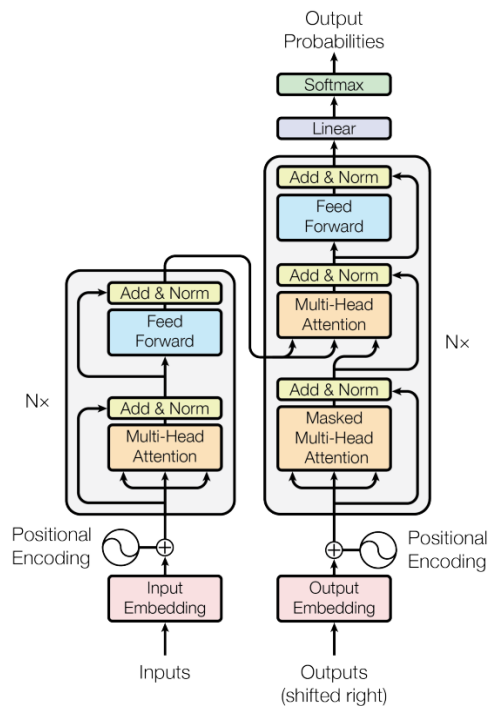


Figure 2.1: Transformer architecture schema [28]

Transformer consists of an Encoder and a Decoder⁹. Figure 2.1 is a schematic representation of the Transformer model. Left part is the Encoder and the right part is the Decoder.

Encoder

The encoder is a stack of multiple identical layers¹⁰, each layer consisting of two sub-layers:

1. Multi-head self-attention layer

⁶This means that it is difficult to parallelize the training process across training examples and smaller batch sizes had to be used.

⁷This means that it is impossible to parallelize training within the training examples.

⁸The hidden state is an intermittent sequence within a sequence-to-sequence model. It is a product of applying the encoder on the input sequence. The output sequence is generated by applying the decoder to the hidden state. Having a hidden state of fixed size means that some information will inherently get lost when a very long input sequence is fed into the model.

⁹Similar to already mentioned RNNs.

¹⁰In the original paper there are 6 of these layers in the encoder stack

125 2. Feed-forward neural network

Multi-head Self-attention layer consists of N so-called "attention heads"¹¹. An attention head is essentially a series of arithmetic operations performed on the input sequence: Let's say we have a sequence of tokens $S = s_1, s_2, \dots, s_n$ the dimension of input tokens is $1 \times d_{\text{model}}$ ¹². An attention head consists of weight matrices¹³ W^Q , W^K and W^V ¹⁴ of the dimension $(d_{\text{model}} \times d_k)$ where $d_k = d_{\text{model}}/N$. Multiplying these weight matrices by token s_i from the input sequence creates Q_i , K_i and V_i vectors respectively of dimension $1 \times d_k$. For each token i , dot product of Q_i and K_j (for all j , $0 < j < n+1$) is calculated: $\text{dot}_{i,1}, \text{dot}_{i,2}, \dots, \text{dot}_{i,n}$. These dot-products are then scaled¹⁵, a softmax layer is applied¹⁶ to all these dot products. Each dot product $\text{dot}_{i,j}$ then multiplies the corresponding V_j to get the series of V vectors: $V_{i,1}, V_{i,2}, \dots, V_{i,n}$. All these V vectors are then element-wise summed up to Z_i (of dimension $1 \times d_k$) which represents the self-attention output for the token i . This is repeated for all n tokens and creates a series of vectors Z_1, Z_2, \dots, Z_n (one for each token) which are then vertically stacked into matrix Z of dimension $n \times d_k$. This matrix Z_h is the output of the single self-attention head h . If we vectorize this process¹⁷, we can express attention as

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

See Figure 2.2:

This process goes on in parallel in every head of the multi-head self-attention layer. Finally, the Z matrices of each self-attention head Z_1, Z_2, \dots, Z_N are horizontally concatenated to produce the output of the whole multi-head self-attention layer L : Z_L (of dimension $n \times N \times d_k$). Multi-head self-attention layer has a weight matrix W_O of dimension $N \times d_k \times d_{\text{model}}$ which gets multiplied by Z_L to produce the final output of the multi-head self-attention layer L .

¹¹In the Original paper, $N = 8$.

¹²In the original paper, $d(\text{model}) = 512$.

¹³These weight matrices are all learnable weights.

¹⁴ Q , K , and V stand for Query, Key and Value.

¹⁵Scaling factor is $1 / \text{square root of } d(k)$. According to the paper, this is done so that the weights are more stable during training.

¹⁶Applying Softmax on a series of numbers means scaling all of them so that they sum up to 1.

¹⁷When we need to do the same arithmetic operations multiple times with different vectors, we can simply concatenate these vectors into matrices and do the operations just once with these matrices.




Figure 2.2: Scaled Dot product attention [28]

Before the output of the multi-head self-attention layer goes into a feed-forward neural network¹⁸, a residual connection¹⁹ [14] and layer normalization²⁰ [2] are applied. Residual
 150 connection and layer normalization are applied after every sub-layer of the Encoder and Decoder.

Decoder

The Decoder has a similar structure as the Encoder with a few notable differences.

Same as the Encoder, Decoder has N layers. Each layer has three sub-layers:

- 155 1. masked self-attention layer
2. encoder-decoder attention layer
3. feed forward layer

In principle, masked self-attention layer works similarly as the multi-head attention

¹⁸This network has two hidden layers with ReLU activation of the first layer and linear activation of the 2nd layer. In the original paper, dimensionality of hidden layers is 2048.

¹⁹Residual connection in this context means adding the input of the layer to the output of that layer before passing to the next layer. This is used to mitigate the vanishing gradient problem and to stabilize the training process.

²⁰Layer normalization is a scaling technique that normalizes the the data across features within each data-sample individually, as opposed to batch normalization which normalizes each feature across multiple data-samples.

layer in the Encoder, however, Q, K, and V come from the already generated output²¹,
 160 and the attention is calculated only for the current and previously generated outputs i.e.
 for an already generated token i , masked self-attention head only calculates Z_1, Z_2, \dots, Z_i .
 Otherwise, the process is the same as in the Encoder multi-head self-attention sub-layer²².
 (See Figure 2.1.)

Encoder-decoder attention layer also works similarly as the multi-head self-attention
 165 layer in the Encoder, however in this layer Q comes from the previous masked self-attention
 layer but K and V come from the Encoder. This allows every position in the Encoder to
 attend all positions in the input. (See Figure 2.1.)

Summary

Even though the Attention mechanism has been around for a while [3], the self-attention
 170 model, as implemented in the Transformer, was a revolutionary step - as it solved the
 issue of long range dependancies²³. This concept enables the model to learn the intrinsic
 structure of the sequence it is presented with and therefore builds a certain level of actual
 "understanding" of the inputs rather than simple "fitting".

Another revolutionary aspect of the Transformer is the positional encoding. Positional
 175 encoding injects information about the absolute and relative position of each token in the
 sequence²⁴ thus allowing the model to learn about the importance of relative positions of
 tokens²⁵.

Finally, and probably most importantly, the way Transformer processes inputs is
 inherently parallelizable - allowing significant scaling benefits to be exploited and very
 180 large models to be built. This is beneficial as larger models can learn more intricate

²¹This makes the Decoder particularly suitable for autoregressive tasks as it predicts tokens sequentially based on which tokens have been predicted before, as opposed to the Encoder which would predict a sequence of tokens which it thinks would have the highest probability of being true overall without regard for the order of the sequence.

²²i.e. Encoder also stacks the outputs of each head, applies the W_o matrix, the residual connection and layer normalization.

²³Earlier attention models had issues relating tokens which were too far apart due to attention being applied only on hidden states (hidden states are a feature of RNN models) rather than on inputs directly.

²⁴this is done by sumating each input token with a positional vector of equal dimension. Value of each number in the positional vector is a sin (or could be cos) function of the position of the token it is being added to as well as of the position inside the positional vector.

²⁵This is most important on the field of NLP (Natural Language Processing) tasks.

patterns in the data and they allow for richer vector representations of the input data²⁶ ²⁷

2.2 Transformer-based models

Many have built on the success of the Transformer and have come up with their own models which retain many concepts from the original Transformer.

NLP Transformer-based models

Among the most famous NLP transformer-based models are the BERT (2018) [10], T5 (2019) [21] which themselves have been built upon many times (2019 alBERT [17], 2019 roBERTa [18], 2019 distilBERT [23], 2020 mT5 [30], 2024 flan-T5 [9]), however, OpenAI's GPT²⁸ series takes the crown as the most famous Transformer-based NLP model(s) by far.

CV (computer vision) Transformer-based models

Among the most famous CV transformer-based models are 2020 ViT [11], 2020 DETR [6], and 2021 swin-Transformer [19].

Speech and Audio processing

²⁶Vector representations in the original Transformer had the dimension 512, whereas some modern transformer-based models support dimensions in the thousands.

²⁷Richer representations of the input data means that the model is able to better understand more nuanced differences between some similar input tokens.

²⁸GPT stands for Generative Pre-trained Transformer

195 **3. RELATED WORK / LITERATURE REVIEW**

3.1 Section

Sub-Section

4. METHODOLOGY

This Research aims to answer 3 main questions:

1. How good are TSFMs on financial time-series data?
2. In which cases do they perform better/worse?
3. How to improve TSFMs performance?

Methodology of this research is designed in a way to try best answer these three questions.

4.1 Time-series prediction evaluation

In order to answer the first question, we need to evaluate time-series predictions. Literature suggests time-series predictions should be evaluated using the traditional regression metrics due to the continuous nature of the target variables. All regression metrics are based on some variation of the prediction error. Prediction error is the difference between the value we predicted and the actual value. From the prediction error, we can derive the following metrics:

- MSE (Mean Squared Error)
- RMSE (Root Mean Square Error)
- MAE (Mean Absolute Error)
- MAPE (Mean Absolute Percentage Error)
- R2

Apart from R2, all these metrics are self-explanatory.

An additional metric that will be used in this research is MDA (Mean Directional Accuracy) and it is time-series specific. Since time-series data is inherently ordered (as opposed to regular regression data - which is unordered), we can measure the direction in

which a time-series is moving at each time-step. Hence we can compare the direction of movement of our predicted time-series against the actual time-series and we can calculate in what percentage of cases did the two time-series move in the same direction (up or down). This measure is the MDA.

225 This is especially important in the field of Finance as in many cases (such as with stock prices) we are not necessarily primarily interested in predicting the actual values, but rather we are interested in predicting the direction in which the price will go next time-step.

4.2 Data

230 This project uses 4 types of financial time-series data:

1. Stock index prices

- (a) S&P 500

- (b) FTSE 100

- (c) DOW Jones

235

- (d) NASDAQ

2. Exchange rates

- (a) USD/GBP

3. Commodity prices

- (a) WTI (Oil)

240

4. Cryptocurrency values

- (a) Bitcoin

The specific time-series were chosen solely based on their prominence in the field of Finance.

Data Pre-processing

245 In the field of Finance, there is a concept of "return". Return of an asset A at a certain point in time T is calculated as

$$(A_T - A_{T-1})/A_{T-1}$$

i.e. the percentage change in the value of that asset between the point in time T and T-1. If we think about each one of the fore-mentioned time-series as assets, we can represent them as time-series of returns rather than their actual values. This is standard practice
250 in Finance as we are not necessarily interested in raw numbers but rather in percentage changes and it is the approach this research will take.

Time-series data characteristics

Time-series data exhibits several key characteristics that are essential for effective analysis and forecasting. One of the primary features is **trend**, which represents the long-term
255 movement in the data, indicating whether the data points generally increase, decrease, or remain constant over time. Another important characteristic is **seasonality**, referring to recurring patterns at regular intervals, often driven by seasonal factors such as months or quarters. In contrast to seasonality, **cyclical patterns** occur over irregular periods and are often influenced by broader economic conditions. Additionally, time-series data is
260 often affected by **noise** or randomness, which represents unpredictable fluctuations that do not follow any discernible pattern. A key property of many time series is **stationarity**, where the statistical properties such as mean and variance remain constant over time. Understanding **autocorrelation**, which measures the correlation of the time series with its lagged values, is crucial for identifying dependencies within the data. **Lag** itself refers to
265 the time step difference between observations that may impact future values. **Volatility**, indicating the degree of variation within the time series, is another significant aspect, especially in financial data where risk management is critical.

Depending on the type and frequency of a financial time-series, it may exhibit different characteristics. For example: Cryptocurrency prices exhibit high volatility

whereas exchange rates exhibit very little volatility or minutely stock index value changes are driven mostly by noise, whereas changes in monthly stock index values are driven mostly by changes in fundamental value of the index. Therefore it is important to run experiments across different types of data and different frequencies to find out how well the models handle different time-series characteristics. Running the experiment with different sizes of training set (context length) and comparing the results, gives us an insight into how well the models handle trend, seasonality and cyclical patterns.

4.3 Experiment Design

Benchmark

Calculating the evaluation metrics for time-series prediction is not sufficient to answer whether a time-series prediction model is good or not. The results of the evaluation need to be put into context and compared against a benchmark in order for us to be able to say whether the results are good or not. Raw numbers don't mean anything.

I have chosen the autoARIMA, simple autoregressor and Meta's Prophet model for the Benchmark models.

Time Series Cross Validation

K-fold Time Series Cross Validation (TSCV) is a statistical method for ensuring statistical significance of a time-series prediction task. In this experiment, I'm using the rolling window version of the TSCV method. Let's say we have chosen the length of the training set to be L , the prediction horizon (how many steps into the future we're predicting) to be H and the time-series of interest has the length L . On the first fold of the TSCV process, we will choose the points $1, 2, \dots, L$ from the time-series to be the "train" set (window) and the points $L+1, L+2, \dots, L+H$ to be the "test" set (window) against which we evaluate our models' predictions. (In the field of finance, we are usually only interested in predicting just the next value in the future therefore, in this research, I will use just $H=1$.) On the next fold of the TSCV, we move the train and test window by one point so we choose points $2, 3, \dots, L+1$ to be the "train" set/window and the points $L+2, L+3, \dots,$

$L+H+1$ to be the "test" set/window. This procedure is repeated K times. By repeating the prediction and evaluation process many times, we hope that the aggregated evaluation of the model's performance that we calculate is statistically representative of that model's

300 actual performance.

5. RESULTS

5.1 Section

Sub-Section

6. DISCUSSION

6.1 Section

Sub-Section

7. CONCLUSION

7.1 Section

Sub-Section

8. APPENDIX

8.1 Brief History of time-series prediction in Finance

When the agricultural revolution took place (circa 10,000 BC), for the first time in history, humans have started producing more food and items than they required just to survive. This gave birth to the concept of "storing" things and eventually trading them for other things which were perceived to be of greater value. When sufficient amount of people started transacting goods in this manner, the concept of "market price" took hold. It did not take long before humans realized that they didn't have to work in the field the whole day to amass wealth, but that they could do it through trade alone. They could buy items from one person at a low price and sell to another at a high price (arbitrage) or buy it now at a low price and hopefully sell later at a high price (asset appreciation). Opportunities to exploit the first option became more rare and less profitable as more and more people started engaging in trade and eliminating these arbitrage opportunities. However, the 2nd option has captured the imagination of traders ever since 10,000BC as no one seemed to be able to tell with certainty whether the price of a good would indeed go up in the future or not. As humans learned to write and keep records, they started noting down the prices of goods that were being traded along with the time when those prices were prevailing. We have records of ancient Babylonians keeping historic records of end-of-month barley, dates, cuscutea, sesame, cardamom and wool prices on clay tablets spanning some 400 years [13]. Investigating these prices, we can see that they follow a random walk pattern with exogenous shocks (such as the death of Alexander the Great) [27] and they seem almost impossible to predict with a naked eye, however, it is not impossible that someone might have tried and unknowingly became the first person in history to work on time-series prediction in the field of Finance.

Over the centuries, methods to solve the problem of time-series prediction in finance seem to have developed universally across many different cultures. Christopher Kurz, a 16th century merchant from Antwerp, thought about the idea of "trend" and "seasonality" in agricultural prices and claimed he could predict prices 20 days in advance. In 18th century Japan, Munehisa Honma - a man dubbed "God of the markets" by his contemporaries, developed a principle stating that when prices become extremely high, they must come

down again which we know today to be the principle of "mean reversion". In the late 19th century, Charles Dow, co-founder of Wall Street Journal and Dow Jones Company²⁹ popularized and coined the term "technical analysis" (a collection of methods that employ math and statistics to predict future prices) which later evolved into quantitative analysis - which is used today in institutions engaged in future price prediction.

8.2 Brief History of modern time-series prediction methods

Statistical models

Although the field of future price prediction has moved far beyond simply considering just the historical prices³⁰, there has been notable work on the general field of univariate³¹ time-series prediction. It is hard to pinpoint the exact start of modern time-series prediction, but I think a good candidate would be Udny Yule's³² 1927 [32] paper which is considered first to have used simple autoregressor model (AR)³³ to predict Wolfer's Sunspot numbers³⁴. Even though Slutsky (1927) [24] and Udny (1921) [31] demonstrated the use of Moving Average as a way to show trend and cyclicity in time-series data, Herman Wold's work (1938)³⁵ [29] indirectly invented the Auto Regressive Moving Average (ARMA) model³⁶. ARMA model is theoretically sound if the time-series being predicted is stationary³⁷, which is not the case with all time-series. In 1970, George Box and Gwilym Jenkins introduced probably the most famous time-series prediction model of all times: ARIMA [4]. ARIMA(p, i, q) model is an ARMA(p, q) model that employs the method of

²⁹Dow Jones Company is the publisher of the prominent DOWJ stock market index

³⁰Quantitative analysis now-adays considers wide palette of data and information

³¹Univariate time-series prediction is a practice of using only the past values of a certain time-series in order to predict its future value(s).

³²Udny studied at UCL under a well-known (and controversial) professor Karl Pearson - father of mathematical statistics and generally one of the most well-known personas in statistics.

³³An AR(p) time-series prediction model predicts future values based on a linear combination of the previous p observations in the series, assuming that the current value depends on its own past values and a stochastic error term.

³⁴Certain time-series data from the field of Astronomy.

³⁵Wold's Decomposition Theorem states that any stationary time series can be decomposed into a deterministic part (AR component) and a stochastic part (MA component)

³⁶The MA(q) part of the ARMA(p, q) model forecasts future values by modeling the current value as a linear combination of the past q error terms (shocks) and a stochastic error term.

³⁷A time series is said to be stationary if its statistical properties, such as mean, variance, and autocorrelation, remain constant over time

differencing³⁸ on the date before training and prediction.

Another direction of innovation in the field of univariate time-series prediction was towards "Exponential Smoothing" (ES). Concept of ES originates from the work of Brown 1956 [5] and has been thoroughly built upon since. Most notable invention was by Hyndman et al. in 2002 [16] - implementing exponential smoothing within a state-space framework, giving us the ETS³⁹ model. Most recent culmination in the ES family of models was by Ivan Svetunkov in 2022 [26] with the CES model⁴⁰. Final honorable mention among the time-series prediction models of "statistical" nature is the "Theta" model. Vassilis Assimakopoulos, and Konstantinos Nikolopoulos (2000) [1] Introduced the Theta model which decomposes a time-series into its' cyclical and trend components which are then forecasted separately and finally combined⁴¹.

Machine Learning (ML) models

As the hardware technology advanced in the 21st century, so did the time-series models as the scientists have finally gotten machines with enough computational power to run more complex algorithms. The most famous of these algorithms being the Artificial Neural Network (ANN). Although the idea (in its primitive form⁴²) dates back all the way to McCulloch and Pitts (1943) [20], ANN first real ancestor was the "Perceptron"⁴³, invented by a psychologist - Frank Rosenblatt in 1958 [22]. The perceptron had ability to take continuous inputs and proper learning algorithm - known even to this day as the "Rosenblatt algorithm". Through the rest of the 20th century, there have been general

³⁸Differencing is a technique in time-series analysis used to transform a non-stationary series into a stationary one by deducting the preceding observation from each observation. The order "i" of the differencing refers to the ammount of times this procedure is applied to a time-series.

³⁹ETS stands for Error Trend Seasonality. ETS is a family of 18 models. ETS model can account for errors being additive or multiplicative, and trend or seasonality being none, additive or multiplicative. ETS models with no multiplicative errors or seasonality have their equivalent withing the ARIMA family of models

⁴⁰CES stands for Complex Exponential Smoothing model. CES method models the time-series as a series of complex numbers where the real part is the prediction and imaginary parts are errors. CES models can also be expressed in state-space form to adress seasonality. CES advantages are that it is flexible and has been empirically shown to perform well.

⁴¹These components are called "theta lines". Theta model can use exponential smoothing or an AR model to predict these two lines. Predictions are finally combined using weighted average where weights could be equal or adjusted to favor either the trend line or the cycle line.

⁴²The original vision of an ANN was based on simple binary inputs with fixed thresholds activation - a far cry from modern ANNs

⁴³Perceptron can be thought of as a single-neuron ANN. ANNs are comprised of one or more neuron layers, each layer containing one or more neurons.

periods of great interest and disinterest in ANNs for many reasons, but ANNs finally came
380 on their own in the 21st century when computers become powerful enough to be able to
train larger versions of these models⁴⁴ in reasonable time⁴⁵. Naturally it was not long
before people realized ANN-type models can also be used for time-series prediction⁴⁶ which
led to many variations of the ANN architecture to be invented. Bayesian regularization
ANN has been used in Finance [12]. Another successful variation of ANN is the Radial
385 Basis Function ANN (RBFANN).

Another

Chen and Guestrin (2016) [7] introduced the Xgboost model which has

Deep Learning models

⁴⁴Size of an ANN matters when dealing with very complex tasks (tasks with many input features).

⁴⁵Another reason for explosion in popularity was the beginning of the internet era and sudden availability of large ammounts of data which are needed for ANN training

⁴⁶It is important to note that univariate time-series prediction ANNs are simply regression ANNs where lags of the data we're trying to predict are explanatory features (inputs) to the model

REFERENCES

- [1] V. Assimakopoulos and K. Nikolopoulos. The theta model: a decomposition approach to forecasting. *International journal of forecasting*, 16(4):521–530, 2000.
- [2] J. Ba. Layer normalization. *arXiv preprint arXiv:1607.06450*, 2016.
- [3] D. Bahdanau. Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*, 2014.
- [4] G. Box and G. Jenkins. *Time Series Analysis: Forecasting and Control*. Holden-Day series in time series analysis and digital processing. Holden-Day, 1970.
- [5] R. G. Brown. *Exponential smoothing for predicting demand*. Little, 1956.
- [6] N. Carion, F. Massa, G. Synnaeve, N. Usunier, A. Kirillov, and S. Zagoruyko. End-to-end object detection with transformers. In *European conference on computer vision*, pages 213–229. Springer, 2020.
- [7] T. Chen and C. Guestrin. Xgboost: A scalable tree boosting system. In *Proceedings of the 22nd acm sigkdd international conference on knowledge discovery and data mining*, pages 785–794, 2016.
- [8] K. Cho. Learning phrase representations using rnn encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078*, 2014.
- [9] H. W. Chung, L. Hou, S. Longpre, B. Zoph, Y. Tay, W. Fedus, Y. Li, X. Wang, M. Dehghani, S. Brahma, et al. Scaling instruction-finetuned language models. *Journal of Machine Learning Research*, 25(70):1–53, 2024.
- [10] J. Devlin. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- [11] A. Dosovitskiy. An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*, 2020.
- [12] R. Gençay and M. Qi. Pricing and hedging derivative securities with neural networks: Bayesian regularization, early stopping, and bagging. *IEEE transactions on neural networks*, 12(4):726–734, 2001.

- [13] J. Hasanhodzic and A. W. Lo. Can hedge-fund returns be replicated?: The linear case. *The Linear Case (August 16, 2006)*, 2006.
- [14] K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.
- [15] S. Hochreiter. Long short-term memory. *Neural Computation MIT-Press*, 1997.
- [16] R. J. Hyndman, A. B. Koehler, R. D. Snyder, and S. Grose. A state space framework for automatic forecasting using exponential smoothing methods. *International Journal of forecasting*, 18(3):439–454, 2002.
- [17] Z. Lan. Albert: A lite bert for self-supervised learning of language representations. *arXiv preprint arXiv:1909.11942*, 2019.
- [18] Y. Liu. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.
- [19] Z. Liu, Y. Lin, Y. Cao, H. Hu, Y. Wei, Z. Zhang, S. Lin, and B. Guo. Swin transformer: Hierarchical vision transformer using shifted windows. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 10012–10022, 2021.
- [20] W. S. McCulloch and W. Pitts. A logical calculus of the ideas immanent in nervous activity. *The bulletin of mathematical biophysics*, 5:115–133, 1943.
- [21] A. Roberts, C. Raffel, K. Lee, M. Matena, N. Shazeer, P. J. Liu, S. Narang, W. Li, and Y. Zhou. Exploring the limits of transfer learning with a unified text-to-text transformer. *Google, Tech. Rep.*, 2019.
- [22] F. Rosenblatt. The perceptron: a probabilistic model for information storage and organization in the brain. *Psychological review*, 65(6):386, 1958.
- [23] V. Sanh. Distilbert, a distilled version of bert: Smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.

- [24] E. SLUTSKY. on a case where the law of large numbers applies to mutually dependent quantities. the extension of a theorem by magoichirô watanabe. *Tohoku Mathematical Journal, First Series*, 28:26–32, 1927.
- 445 [25] I. Sutskever. Sequence to sequence learning with neural networks. *arXiv preprint arXiv:1409.3215*, 2014.
- [26] I. Svetunkov, N. Kourentzes, and J. K. Ord. Complex exponential smoothing. *Naval Research Logistics (NRL)*, 69(8):1108–1123, 2022.
- [27] P. Temin. Price behavior in ancient babylon. *Explorations in Economic History*,
450 39(1):46–60, 2002.
- [28] A. Vaswani. Attention is all you need. *Advances in Neural Information Processing Systems*, 2017.
- [29] H. Wold. *A study in the analysis of stationary time series*. PhD thesis, Almqvist & Wiksell, 1938.
- 455 [30] L. Xue. mt5: A massively multilingual pre-trained text-to-text transformer. *arXiv preprint arXiv:2010.11934*, 2020.
- [31] G. U. Yule. On the time-correlation problem, with especial reference to the variate-difference correlation method. *Journal of the Royal Statistical Society*, 84(4):497–537, 1921.
- 460 [32] G. U. Yule. Vii. on a method of investigating periodicities disturbed series, with special reference to wolfer’s sunspot numbers. *Philosophical Transactions of the Royal Society of London. Series A, Containing Papers of a Mathematical or Physical Character*, 226(636-646):267–298, 1927.